

Fig. 1. Average number of iterations in 1D as a function of L .

1. Programming a D-dimensional random walk

The general approach was to use D indices stored in an array to track the position of the particle within each dimension. By first generating a random number to pick the dimension and then a second to move forwards or backwards I was able to simulate a multi-dimensional random walk. To track the position, each coordinate was converted into a unique index J using,

$$J = \sum_{j=0}^{D-1} x_j L^j, \quad (1)$$

where x_j is the index in a given dimension. The benefit of flattening the array, as opposed to tracking the position in a D -dimensional array, is that it's much easier to scale to higher dimensions while maintaining manageable dynamic memory allocation. Information regarding the particles position within the grid is stored inside a structure. Memory is then dynamically allocated for the position and tracker arrays, that get resized as either L or D changes. At each step of the random walk, the current coordinate is converted into J and the tracker array is checked at that index. If its value is 0 then it, along with a separate *count* variable, get incremented by 1. This process continues until the value of *count* is equal to the total number, $N = L^D$, of positions within the grid, signifying all positions have been visited.

2. 1D random walk

2.1. Seeding random()

random() will use a linear congruential algorithm,

$$X_{n+1} = (aX_n + c) \bmod m, \quad (2)$$

or a similar formula, so the same seed X_0 will naturally produce the same set of pseudo-random numbers. In later code I have chosen to use the current time as the seed.

2.2. Average iterations for $L = 10$

The average number of hops converges to 46.00 after 10^7 iterations

2.3. Average iterations with varying L

Figure 1 has been calculated using 10^5 iterations for the first 100 values of L and 10^3 for the remaining 900. From this we find the power law relationship describing the average number of iterations,

$$\bar{I} = (0.53 \pm 0.011) L^{(1.99 \pm 0.0030)}, \quad (3)$$




Fig. 2. $\bar{I} = (7.86 \pm 0.26) N^{(1.073 \pm 0.002)}$ for $0 < L < 160$, each at 100 iterations, in 3D.




Fig. 3. Decreased scaling of \bar{I} with N in 1D (red), 2D (blue), 3D (green), 4D (yellow).

where the error derives from SciPy CURVE-FIT. To examine the accuracy of this equation, I simulated $L = 2000$ for 10^4 iterations, which returns $\bar{I} = 2.01 \times 10^6$, very similar to $\bar{I} = 1.97 \times 10^6$ obtained from this formula, suggesting suitable extrapolation. From the central limit theorem, given a sufficiently large number of iterations, one expects a normal distribution to appropriately describe the probability that the random walk is a certain distance from the start. This suggests that as L increases the probability of the particle ending up at the furthest position falls off in a nonlinear fashion. Therefore, a power law relationship appropriately describes the average number of iterations, $\bar{I}(L)$, to visit all positions.

3. Higher dimensional random walk

As $N = L^D$, in higher dimensions we no longer have a 1:1 mapping between N and L . Instead for the same N the grid is far more connected. A particles position is described by D indices, therefore one can imagine these as D independent particles each confined to a given dimension. In a random walk these 'particles' will move on average every D goes, each with a normal probability distribution about the origin. For a given N , the number of positions in 1D within the grid is $L = N^{1/D}$; thus, each particle is far more likely to explore the extremes of the space. Further, as D increases, the number of directions available to the particle also grows as $2 \times D$, meaning the particle has a lower probability of revisiting sites, and hence for a given N a fewer number of iterations to explore the available space.




Fig. 4. $\bar{I}(D)$ for the first 13 dimensions, at $L = 3$ and 100 iterations.

Figure 2 shows the increased linearity of $\bar{I}(N)$ in 3D due to these effects, and Figure 3 shows $\bar{I}(N)$ becoming flatter as D varies from 1 to 4. Increasing D further, while holding $L = 3$ fixed, naturally reveals the relationship of $\bar{I}(D)$ depicted in Figure 4. This power-law increase follows from $N = L^D$.

4. Random walks on a network

4.1. Computational approach

Much of the same code can be reused here with the main changes being to the indexing and random number generation. We can label each component, i , from 1 to A , and nodes, j , within each component from 1 (the node connected to $i - 1$) to B (the node connected to $i + 1$). If $j \neq 1, B$ then the particle can move to any other maximumly connected node j within i , with probability $1/(B - 1)$. Thus, we can use a simple acceptance-rejection algorithm to generate a number between 1 and B until we get a number that is not the current position. For a reasonable size of B this has a favourable rejection probability of only $1/B$. If instead $j = 1, B$ then the particle may also move components. We can do almost the same as above but generate a random number x starting at 0. If $x > 0$ then we move to node $j = x$ within i , but if $x = 0$ then we move to the next component and j changes to 1, if B originally, or vice versa.

4.2. Limitations

If A or $B = 1$, then the simulation results are inaccurate, and while these could be solved with an *if* statement to catch these specific cases, I have chosen not to as these conditions are unnecessary for the reasons outlined below, and would clutter the code. If $A = 1$, the periodic boundary condition creates two edges between nodes 1 and B . However, this particular case can be solved analytically as \bar{I} is equivalent to the expected number of outcomes needed to generate every random number between 1 and B , which can be solved easily using

$$\bar{I} = E(B) = B \times H_n, \quad (4)$$

where H_n is the sum of the reciprocals of the first B integers. If instead B equals 1, a 1D random walk is returned. These limitations are greatly outweighed by the memory and speed advantages of not having a matrix store which nodes are connected to one another.

4.3. Convergence to a 1D random walk

Setting $B = 2$ recovers an even 1D random walk, and with $A = 5$ we find $\bar{I} = 46.00$, the same as for $L = 10$.




Fig. 5. $\bar{I}(A)$ (red) at fixed $B = 4$, and $\bar{I}(B)$ (blue) at fixed $A = 5$, both averaged over 10^5 iterations at each point.

4.4. Fixed and dynamic A, B

For fixed $A = 5, B = 4$ we find $\bar{I} = 167.54$, averaged for 10^7 iterations. Figure 5 shows A and B varying between 1 and 100 while the other variable is held fixed. Interestingly $\bar{I}(A) \propto A^{(1.99 \pm 0.004)}$ scales the same as for the 1D walk ($\propto L^{1.99}$). This remains true for increasing values of B too, but with an enlarging proportionality constant. This suggests that the network random walk is much like that of the 1D walk, except scaled by a larger pre-factor, as the particle can now get stuck in a component of B nodes at each position. As B increases the particle spends a greater period of time within each component, but scales at a lesser rate (for equated A, B) than $\bar{I}(A)$. This may be because, from Equation 4, $E(B) \propto B^{1.11}$, reducing the scaling factor.

4.5. Multi-particle diffusion

Physical diffusion problems involve many more than one particle, so to simulate this I created an array of ‘particle’ structures that can move independently. For n particles, one can then track the average number of times each component is visited for a given number of iterations. Figure 6 illustrates these results, depicting a normal distribution of the particles about the start position as discussed in section 3, offering a promising validation of the theory.




Fig. 6. Normalised distribution, centred at component 51 (start position), depicting the number of times a component, i , is visited for 5 particles and 10^4 iterations at $A = 100$.